

- Detailed description of your cloud benchmarking methodology, including any scripts or other code

1) Querying the list of your VMs

For this scenario we executed Nova_list_servers script which takes time and concurrency as parameters. For obtaining a better average we gave the following to time and concurrency:

Time: 3

Concurrency: 10

To calibrate the results we checked it against the OpenStack command “openstack server list” and got equivalent performance in ms.

Script: NovaServers.list_servers.JSON

```
{
  "NovaServers.list_servers": [
    {
      "args": {
        "detailed": true
      },
      "context": {},
      "hooks": [],
      "runner": {
        "concurrency": 3,
        "times": 10,
        "type": "constant"
      },
      "sla": {}
    }
  ]
}
```

2) Creation Time of VMs (including the boot time)

For cloud benchmarking of creation/boot scenario we have used “NovaServers.boot_server” script as example. This benchmark scenario allows to test with different parameters, image and flavors.

The fields we used includes

Runner: The type selected here is ‘constant’ since the scenario required on different time intervals at particular times of a day and we had limited number of instances that could be initialized.

Time: The number of times iteration has to be executed for benchmarking. We have kept this value as **1** to allow concurrent execution of boot scripts keeping in view limited quota

Concurrency: The number of times the script executes in parallel. We gave it value **3**.

Checking concurrent initiation/boot of our VMs is the most important factor in cloud benchmarking. This was tested in three different time slots. Apart from some observed downtime during night time, the difference in performance results is negligible.

Script: NovaServers.boot_server.JSON

```
{
  "NovaServers.boot_server": [
    {
      "args": {
        "flavor": {
          "name": "Cloud Computing"
        },
        "image": {
          "name": "ubuntu-16.04"
        }
      },
      "context": {},
      "hooks": [],
      "runner": {
        "concurrency": 3,
        "times": 3,
        "type": "constant"
      },
      "sla": {}
    }
  ]
}
```

- Benchmarking results of six different scenarios and time slots, including plots and interpretation of results

NovaServers.list_servers													
S.No	Load Duration	Full Duration	Iterations	Failures	Started at	Min (sec)	Median (sec)	90%ile (sec)	95%ile (sec)	Max (sec)	Avg (sec)	Success	Count
1	5.424s	9.152s	10	0	2017-07-02T00:33:05	1.187	1.347	1.729	1.929	2.129	1.455	1	10
2	3.921	6.893	10	0	2017-07-01T11:00:02	0.815	1.006	1.221	1.324	1.324	1.036	1	10
3	5.512	8.899s	10	0	2017-07-01T17:27:05	1.003	1.247	1.689	1.265	2.177	1.337	1	10
NovaServers.boot_server													
S.No	Load Duration	Full Duration	Iterations	Failures	Started at	Min (sec)	Median (sec)	90%ile (sec)	95%ile (sec)	Max (sec)	Avg (sec)	Success	Count
1	31.621 s	82.424 s	3	0	2017-07-01T23:36:01	29.289	32.393	32.572	32.595	32.617	31.433	1	3
2	14.216	30.423	3	0	2017-07-02T10:30:02	14.216	15.045	15.052	15.053	15.053	14.772	1	3
3	25.458	74.569	3	0	2017-07-01T16:23:05	24.369	25.487	26.874	26.911	26.915	25.621	1	3

From this we can visibly see that during the morning and evening openstack work better than in the night.

- Commented listing of commands you executed for Task 2

```
//Task-1

//Loading environment variables
source openrc

//Creating rally openstack deployment

//Task-2

//Loading environment variables
source openrc

//Instantiate server.yaml file to start instance passing all the arguments
as parameters
openstack stack create --template server.yaml/ --parameter
"name=Server;image=ubuntu-16.04;flavor=Cloud Computing;key_pair=Cloud-
Computing;network=tu-internal;zone=Cloud Computing 2017" Mystack

//checking the status of the instance
openstack stack list

//Creating the floating ip for assigning it to the server.
openstack floating ip create tu-internal

//Assigning the floating ip to the server
openstack server add floating ip Server 10.200.1.154

//Adding ICMP and SSH rule so that the created VM can be pinged and
connect to SSH
openstack security group rule create default --protocol tcp --dst-port
22:22 --remote-ip 0.0.0.0/0
openstack security group rule create --protocol icmp default

//Ping the VM
ping 10.200.1.154

//Connect the VM through SSH
sudo ssh -i Cloud-Computing ubuntu@10.200.1.154

//Check the internet connectivity of the VM
ping google.com

//Delete the stack
openstack stack delete Mystack
```

```
//Check the Stack have been deleted
openstack stack list
```

- The contents of your server-landscape.yml file

heat_template_version: 2015-10-15

description: Create a new neutron network plus a router to the public

network, and for deploying one frontend servers and two backend server into the new network. The template also

assigns floating IP addresses to the frontend server so they are routable from the public network. and also assigning custom security groups to the front end server.

parameters:

name:

type: string

label: Name of the VM

key_pair:

type: string

label: Key Pair

constraints:

- custom_constraint: nova.keypair

flavor:

type: string

label: Flavor

constraints:

- custom_constraint: nova.flavor

image:

type: string

label: Image Name

constraints:

- custom_constraint: glance.image

network:

type: string

label: Network

constraints:

- # - custom_constraint: neutron.network

zone:

type: string

label: Availability Zone

default: Default

security_groups:

type: comma_delimited_list

label: Security Group(s)

default: "default"

cidr_private:

type: string

label: cidr private

gateway_ip:

type: string

label: gateway_ip

networkcloud:

type: string

label: networkcloud

public_network:

type: string

label: public_network

resources:

This port is a separate resource used to assign the security groups

to the VM. Can also be used to attach a OS::Neutron::FloatingIP to the VM.

instance:

type: server.yaml

properties:

name: { get_param: name }

key_pair: { get_param: key_pair }

image: { get_param: image }

flavor: { get_param: flavor }

zone: { get_param: zone }

network: { get_attr: [private_subnet, network_id]}

#subnet: {get_resource: private_subnet}

security_groups:

- default

- {get_resource : secgroup}

my_backend_group:

type: OS::Heat::ResourceGroup

properties:

count: 2

resource_def:

type: server.yaml

properties:

name: backend-%index%

key_pair: { get_param: key_pair }

image: { get_param: image }

flavor: { get_param: flavor }

zone: { get_param: zone }

network: { get_attr: [private_subnet, network_id]}

#subnet: {get_resource: private_subnet}

secgroup:

type: OS::Neutron::SecurityGroup

properties:

rules:

- protocol: tcp

remote_ip_prefix: 0.0.0.0/0

port_range_min: 80

port_range_max: 80

- protocol: icmp

- protocol: tcp

port_range_min: 22

port_range_max: 22

private_net:

type: OS::Neutron::Net

properties:

name: {get_param: networkcloud}

private_subnet:

type: OS::Neutron::Subnet

properties:

network_id : { get_resource : private_net }

cidr: {get_param: cidr_private}

gateway_ip: {get_param: gateway_ip}

router:

type: OS::Neutron::Router

properties:

external_gateway_info:

network: {get_param : public_network}

router_interface:

type: OS::Neutron::RouterInterface

properties:

router_id : {get_resource: router}

subnet_id : {get_resource: private_subnet}

floating_ip:

type: OS::Neutron::FloatingIP

properties:

floating_network: { get_param: public_network }

floating_ip_assoc:

type: OS::Neutron::FloatingIPAssociation

properties:

floatingip_id: { get_resource: floating_ip }

port_id: { get_attr: [instance,port] }

outputs:

test_out:

description: Value of server id

value: {get_attr: [instance,server]}

details:

description: Details of server

value: {get_attr: [instance]}

floating ip:

description: Floating IP of the instance

value: { get_attr: [floating_ip, floating_ip_address] }

- Commented listing of commands you executed to test your advanced Heat template

```
//Task-3

//Loading environment variables
source openrc

//Instantiate server-landscape.yaml file to start instance passing all the
arguments as parameters
openstack stack create --template A2/server-landscape.yaml/ --parameter
"name=frontend;image=ubuntu-16.04;flavor=Cloud Computing;key_pair=Cloud-
Computing;cidr_private=10.12.1.0/24;gateway_ip=10.12.1.1;zone=Cloud
Computing 2017;public_network=tu-internal;networkcloud=Cloud-network"
Mystack

//checking the status of the instance
openstack stack list

//Extracting the floating ip of the Front-end server
openstack stack show mystack

//Pinging the VM
ping 10.200.1.163

//Transfer the private key for accessing back-end server from the frontend
servers
sudo scp -i Cloud-Computing Cloud-Computing ubuntu@10.200.1.163:

//Connect the frontend through SSH
sudo ssh -i Cloud-Computing ubuntu@10.200.1.163

//Check the internet connectivity of the VM
ping google.com

//Connect the backend server through SSH from frontend server
sudo ssh -i Cloud-Computing ubuntu@10.12.1.5

//Check the internet connectivity of the VM
ping google.

//Exit the backend server
exit

//Connect the backend server through SSH from frontend server
sudo ssh -i Cloud-Computing ubuntu@10.12.1.4

//Check the internet connectivity of the VM
ping google.com
```

```
//Exit the backend server  
exit
```

```
//Deleting the stack  
openstack stack delete Mystack
```

```
//Check for the Stack have been deleted  
openstack stack list
```